

Preparing data using data generator

```
In [6]: train_datagen = ImageDataGenerator(
        zoom_range = 0.2,
        shear_range = 0.2,
        horizontal_flip=True,
        rescale = 1./255
    )

train_data = train_datagen.flow_from_directory(directory= "/content/train",
                                              target_size=(224,224),
                                              batch_size=32,
                                              )

train_data.class_indices
```

Found 28709 images belonging to 7 classes.

```
Out[6]: {'angry': 0,
        'disgust': 1,
        'fear': 2,
        'happy': 3,
        'neutral': 4,
        'sad': 5,
        'surprise': 6}
```

```
In [7]: val_datagen = ImageDataGenerator(rescale = 1./255 )

val_data = val_datagen.flow_from_directory(directory= "/content/test",
                                           target_size=(224,224),
                                           batch_size=32,
                                           )
```

Found 7178 images belonging to 7 classes.

Visualising the data fed

```
In [ ]:
```

```
In [8]: # to visualize the images in the traing data denerator

t_img , label = train_data.next()

#-----
# function when called will prot the images
def plotImages(img_arr, label):
    """
    input  :- images array
    output :- plots the images
    """
    count = 0
    for im, l in zip(img_arr,label) :
        plt.imshow(im)
        plt.title(im.shape)
        plt.axis = False
        plt.show()

        count += 1
        if count == 10:
            break

#-----
# function call to plot the images
plotImages(t_img, label)
```



Early Stopping and Check points

```
In [9]: ## having early stopping and model check point

from keras.callbacks import ModelCheckpoint, EarlyStopping

# early stopping
es = EarlyStopping(monitor='val_accuracy', min_delta= 0.01 , patience= 5, verbose= 1, mode='auto')

# model check point
mc = ModelCheckpoint(filepath="best_model.h5", monitor= 'val_accuracy', verbose= 1, save_best_only= True, mode = 'auto')

# puting call back in a List
call_back = [es, mc]
```

```
In [10]: hist = model.fit_generator(train_data,
                                steps_per_epoch= 10,
                                epochs= 30,
                                validation_data= val_data,
                                validation_steps= 8,
                                callbacks=[es,mc])

<ipython-input-10-7dbb6b2d1ed1>:1: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
    hist = model.fit_generator(train_data,

Epoch 1/30
10/10 [=====] - ETA: 0s - loss: 17.3073 - accuracy: 0.2250
Epoch 1: val_accuracy improved from -inf to 0.19531, saving model to best_model.h5

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
    saving_api.save_model(


10/10 [=====] - 10s 450ms/step - loss: 17.3073 - accuracy: 0.2250 - val_loss: 14.0523 - val_accuracy: 0.1953
Epoch 2/30
10/10 [=====] - ETA: 0s - loss: 14.6760 - accuracy: 0.2719
Epoch 2: val_accuracy improved from 0.19531 to 0.30078, saving model to best_model.h5
10/10 [=====] - 8s 896ms/step - loss: 14.6760 - accuracy: 0.2719 - val_loss: 9.5517 - val_accuracy: 0.3008
Epoch 3/30
10/10 [=====] - ETA: 0s - loss: 9.4273 - accuracy: 0.3625
Epoch 3: val_accuracy improved from 0.30078 to 0.32031, saving model to best_model.h5
10/10 [=====] - 4s 427ms/step - loss: 9.4273 - accuracy: 0.3625 - val_loss: 9.3336 - val_accuracy: 0.3203
Epoch 4/30
10/10 [=====] - ETA: 0s - loss: 8.1260 - accuracy: 0.3594
Epoch 4: val_accuracy improved from 0.32031 to 0.41406, saving model to best_model.h5
10/10 [=====] - 6s 574ms/step - loss: 8.1260 - accuracy: 0.3594 - val_loss: 5.8851 - val_accuracy: 0.4141
Epoch 5/30
10/10 [=====] - ETA: 0s - loss: 7.4087 - accuracy: 0.3625
Epoch 5: val_accuracy did not improve from 0.41406
10/10 [=====] - 4s 408ms/step - loss: 7.4087 - accuracy: 0.3625 - val_loss: 6.8454 - val_accuracy: 0.3438
Epoch 6/30
10/10 [=====] - ETA: 0s - loss: 6.0202 - accuracy: 0.4156
Epoch 6: val_accuracy did not improve from 0.41406
10/10 [=====] - 5s 430ms/step - loss: 6.0202 - accuracy: 0.4156 - val_loss: 6.7367 - val_accuracy: 0.3828
Epoch 7/30
10/10 [=====] - ETA: 0s - loss: 5.7904 - accuracy: 0.4563
Epoch 7: val_accuracy did not improve from 0.41406
10/10 [=====] - 4s 413ms/step - loss: 5.7904 - accuracy: 0.4563 - val_loss: 5.9508 - val_accuracy: 0.3828
Epoch 8/30
10/10 [=====] - ETA: 0s - loss: 4.9099 - accuracy: 0.4563
Epoch 8: val_accuracy improved from 0.41406 to 0.42969, saving model to best_model.h5
10/10 [=====] - 5s 566ms/step - loss: 4.9099 - accuracy: 0.4563 - val_loss: 5.0625 - val_accuracy: 0.4297
Epoch 9/30
10/10 [=====] - ETA: 0s - loss: 5.3911 - accuracy: 0.4313
Epoch 9: val_accuracy improved from 0.42969 to 0.44531, saving model to best_model.h5
10/10 [=====] - 4s 425ms/step - loss: 5.3911 - accuracy: 0.4313 - val_loss: 4.3755 - val_accuracy: 0.4453
Epoch 10/30
10/10 [=====] - ETA: 0s - loss: 4.8029 - accuracy: 0.4375
Epoch 10: val_accuracy improved from 0.44531 to 0.46094, saving model to best_model.h5
10/10 [=====] - 5s 528ms/step - loss: 4.8029 - accuracy: 0.4375 - val_loss: 5.0747 - val_accuracy: 0.4609
Epoch 11/30
10/10 [=====] - ETA: 0s - loss: 5.0435 - accuracy: 0.4313
Epoch 11: val_accuracy did not improve from 0.46094
10/10 [=====] - 4s 400ms/step - loss: 5.0435 - accuracy: 0.4313 - val_loss: 4.6358 - val_accuracy: 0.3828
Epoch 12/30
10/10 [=====] - ETA: 0s - loss: 5.1884 - accuracy: 0.4531
Epoch 12: val_accuracy did not improve from 0.46094
10/10 [=====] - 5s 494ms/step - loss: 5.1884 - accuracy: 0.4531 - val_loss: 5.3386 - val_accuracy: 0.3633
Epoch 13/30
10/10 [=====] - ETA: 0s - loss: 5.6525 - accuracy: 0.4250
Epoch 13: val_accuracy improved from 0.46094 to 0.48828, saving model to best_model.h5
10/10 [=====] - 5s 427ms/step - loss: 5.6525 - accuracy: 0.4250 - val_loss: 4.5562 - val_accuracy: 0.4883
Epoch 14/30
10/10 [=====] - ETA: 0s - loss: 5.4321 - accuracy: 0.4187
Epoch 14: val_accuracy did not improve from 0.48828
10/10 [=====] - 4s 409ms/step - loss: 5.4321 - accuracy: 0.4187 - val_loss: 6.0207 - val_accuracy: 0.4492
Epoch 15/30
10/10 [=====] - ETA: 0s - loss: 6.2157 - accuracy: 0.4300
Epoch 15: val_accuracy did not improve from 0.48828
10/10 [=====] - 5s 553ms/step - loss: 6.2157 - accuracy: 0.4300 - val_loss: 7.7263 - val_accuracy: 0.3516
Epoch 16/30
10/10 [=====] - ETA: 0s - loss: 7.4774 - accuracy: 0.4094
Epoch 16: val_accuracy did not improve from 0.48828
10/10 [=====] - 4s 411ms/step - loss: 7.4774 - accuracy: 0.4094 - val_loss: 8.1603 - val_accuracy: 0.3672
Epoch 17/30
10/10 [=====] - ETA: 0s - loss: 5.7995 - accuracy: 0.4469
Epoch 17: val_accuracy did not improve from 0.48828
10/10 [=====] - 5s 437ms/step - loss: 5.7995 - accuracy: 0.4469 - val_loss: 6.1236 - val_accuracy: 0.4648
Epoch 18/30
10/10 [=====] - ETA: 0s - loss: 5.4460 - accuracy: 0.4812
Epoch 18: val_accuracy did not improve from 0.48828
10/10 [=====] - 4s 400ms/step - loss: 5.4460 - accuracy: 0.4812 - val_loss: 5.6344 - val_accuracy: 0.4453
Epoch 18: early stopping
```

```
In [11]: # Loading the best fit model
from keras.models import load_model
model = load_model("/content/best_model.h5")
```


```
In [12]: h = hist.history
h.keys()
```

Out[12]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```
In [13]: plt.plot(h['accuracy'])
plt.plot(h['val_accuracy'] , c = "red")
plt.title("acc vs v-acc")
plt.show()
```



```
In [14]: plt.plot(h['loss'])
plt.plot(h['val_loss'] , c = "red")
plt.title("loss vs v-loss")
plt.show()
```



```
In [15]: # just to map o/p values
op = dict(zip( train_data.class_indices.values(), train_data.class_indices.keys()))
```

```
In [16]: # path for the image to see if it predics correct class

path = "/content/test/angry/PrivateTest_1054527.jpg"
img = load_img(path, target_size=(224,224) )

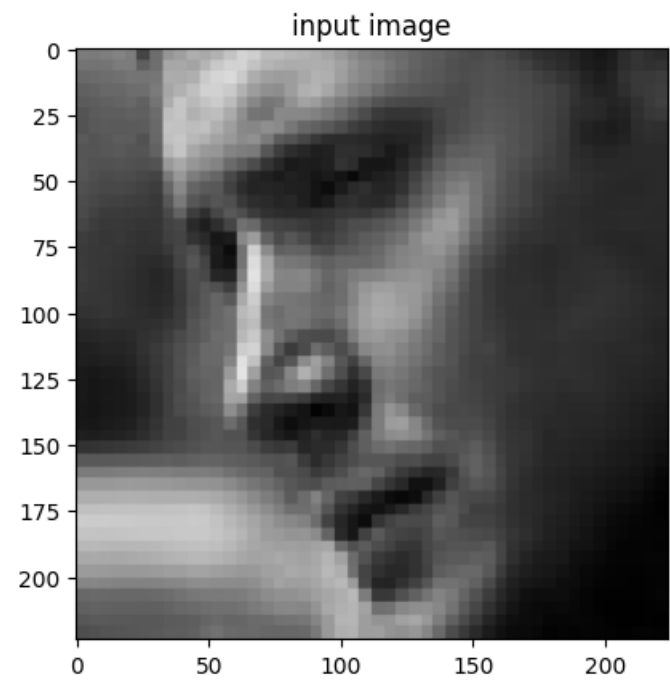
i = img_to_array(img)/255
input_arr = np.array([i])
input_arr.shape

pred = np.argmax(model.predict(input_arr))

print(f" the image is of {op[pred]}")

# to display the image
plt.imshow(input_arr[0])
plt.title("input image")
plt.show()
```

1/1 [=====] - 1s 1s/step
the image is of sad



After testing extensively, it's clear that our model does a good job predicting emotions in images.

In []: Submitted by- Shweta Kanungo

In []: